

```
1 import numpy as np
2 import pandas as pd
3 import altair as alt
4 import streamlit as st
5 import yfinance as yf
6 import matplotlib.pyplot as plt
7
8 from sklearn.preprocessing import MinMaxScaler
9 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
10
11 from tensorflow.keras.models import Sequential
12 from tensorflow.keras.layers import Dense, LSTM, Dropout, Input
13 from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
14
15 # -----
16 # Helpers
17 # -----
18
19 def make_xy(df_scaled, look_back):
20     X, y = [], []
21     for i in range(len(df_scaled) - look_back):
22         X.append(df_scaled[i : i + look_back, 0])
23         y.append(df_scaled[i + look_back, 0])
24     X = np.array(X).reshape(-1, look_back, 1)
25     y = np.array(y).reshape(-1, 1)
26     return X, y
27
28
29 def build_model(look_back, units, dense, d1, d2):
30     model = Sequential([
31         Input(shape=(look_back, 1)),
32         LSTM(units, return_sequences=True),
33         Dropout(d1),
34         LSTM(units, return_sequences=False),
35         Dropout(d2),
36         Dense(dense, activation="relu"),
37         Dense(1),
38     ])
39     model.compile(optimizer="adam", loss="mse", metrics=["RootMeanSquaredError"])
40     return model
41
42
43 def series_metrics(y_true, y_pred):
44     y_true = np.asarray(y_true).ravel()
45     y_pred = np.asarray(y_pred).ravel()
46
47     mae = float(mean_absolute_error(y_true, y_pred))
48     rmse = float(np.sqrt(mean_squared_error(y_true, y_pred)))
49     r2 = float(r2_score(y_true, y_pred))
50
51     return {
52         "MAE_$": mae,
53         "RMSE_$": rmse,
```

```

54         "R2": r2,
55     }
56
57
58     # -----
59     # UI
60     # -----
61
62     st.set_page_config(page_title="Stock Closing Price Forecast (LSTM)", layout="wide")
63     st.title("📈 Stock Closing Price Forecast – LSTM")
64     st.caption("Data source: Yahoo Finance (3 years history by default)")
65
66     with st.sidebar:
67         st.header("1) Data")
68         period = st.selectbox("Download period", ["1y", "2y", "3y", "5y", "10y"], index=2)
69         ticker = st.selectbox("Ticker", ["AAPL", "AMZN", "MSFT", "GOOGL", "NVDA", "TSLA",
70 "META"], index=1)
71
72         st.header("2) Window & Horizon")
73         look_back = st.number_input("Look-back window (days)", min_value=10, max_value=200,
74 value=30, step=5)
75         horizon = st.number_input("Forecast horizon (days)", min_value=1, max_value=60,
76 value=14, step=1)
77
78         units = 64
79         dense = 32
80         d1 = 0.2
81         d2 = 0.1
82         epochs = 50
83
84         test_frac = 0.2
85
86     try:
87         df = yf.Ticker(ticker).history(period=period)
88     except Exception as e:
89         st.error(f"yfinance error: {e}")
90         st.stop()
91
92     if df is None or df.empty:
93         st.error("No data returned. Try another period or check ticker symbol.")
94         st.stop()
95
96     # Ensure Date index and keep Close
97     df = df.reset_index().set_index("Date").sort_index()
98     prices = df[["Close"]].dropna().copy()
99
100     # normalize timezone -> naive (no tz)
101     try:
102         prices.index = prices.index.tz_localize(None)
103     except TypeError:
104         pass
105
106     min_d = prices.index.min().date()
107     max_d = prices.index.max().date()

```

```
105 c1, c2 = st.columns(2)
106
107 with c1:
108     start_date = st.date_input("Start date", value=min_d, min_value=min_d,
109     max_value=max_d)
110 with c2:
111     end_date = st.date_input("End date", value=max_d, min_value=min_d, max_value=max_d)
112
113 mask = (prices.index >= pd.Timestamp(start_date)) & (prices.index <=
114 pd.Timestamp(end_date))
115 series = prices.loc[mask].copy()
116
117 # 50 is way to just say buffer to say
118 if len(series) < look_back + 50:
119     st.warning("Selected window looks short for LSTM. Consider more history.")
120
121 st.subheader(f"{ticker} Stock Data preview")
122 st.dataframe(series.tail(10))
123
124 # Chart raw close
125 raw_chart = (
126     alt.Chart(series.reset_index())
127     .mark_line()
128     .encode(
129         x=alt.X("Date:T", axis=alt.Axis(format="%Y-%m-%d", title="Date")),
130         y=alt.Y("Close:Q", title="Close ($)"),
131         tooltip=[
132             alt.Tooltip("Date:T", format="%Y-%m-%d"),
133             alt.Tooltip("Close:Q", format="$.2f"),
134         ],
135     )
136     .properties(height=250)
137 )
138 st.altair_chart(raw_chart, use_container_width=True)
139
140 values = series.values.astype(float)
141 N = len(values)
142
143 if N <= look_back + 1:
144     st.error("Not enough data for chosen look-back.")
145     st.stop()
146
147 # Time-based split
148 train_size = int(N * (1 - float(test_frac)))
149 train_vals = values[:train_size] # 0... train_size
150 test_vals = values[train_size:] # train_size...end
151
152 # Scaling (fit on train only)
153 scaler = MinMaxScaler()
154 train_scaled = scaler.fit_transform(train_vals)
155 test_scaled = scaler.transform(test_vals)
156
157 # Windows
158 X_train, y_train = make_xy(train_scaled, look_back)
```

```

157 X_test, y_test = make_xy(test_scaled, look_back)
158
159 if len(X_train) == 0 or len(X_test) == 0:
160     st.error("Look-back too large for the selected split; reduce look-back or expand date
range.")
161     st.stop()
162
163 # Train & Forecast
164 if st.button("🚀 Train & Forecast", use_container_width=True):
165     with st.spinner("Training model..."):
166         model = build_model(look_back, units, dense, d1, d2)
167         callbacks = [
168             EarlyStopping(patience = 10, restore_best_weights = True),
169             ReduceLROnPlateau(),
170         ]
171         model.fit(
172             X_train,
173             y_train,
174             batch_size=32,
175             validation_data=(X_test, y_test),
176             epochs=int(epochs),
177             callbacks=callbacks
178         )
179
180     # Predict test
181     y_pred_scaled = model.predict(X_test) ## (n_samples,1)
182     y_pred = scaler.inverse_transform(y_pred_scaled).ravel() ## flatten to 1D array
183     y_true = scaler.inverse_transform(y_test).ravel()
184
185     # Index alignment
186     test_index = series.index[train_size + look_back:]
187     test_df = pd.DataFrame({"Actual_$": y_true, "Pred_$": y_pred}, index=test_index)
188
189     # Metrics
190     met = series_metrics(y_true, y_pred)
191
192     st.subheader("Evaluation Summary")
193     m1, m2, m3 = st.columns(3)
194     m1.metric("MAE ($)", f"{met['MAE_$']:.2f}")
195     m2.metric("RMSE ($)", f"{met['RMSE_$']:.2f}")
196     m3.metric("R²", f"{met['R2']:.3f}")
197
198     # Plot test
199     st.markdown("### Test: Actual vs Predicted")
200
201     test_chart = (
202         alt.Chart(test_df.reset_index())
203         #transform_fold is the key: it turns two columns into a single "series" field so
Altair can draw two colored lines from one tidy table.
204         .transform_fold(["Actual_$", "Pred_$"], as_=["variable", "value"])
205         .mark_line()
206         .encode(
207             x=alt.X("Date:T", axis=alt.Axis(format="%Y-%m-%d", title="Date")),
208             y=alt.Y("value:Q", title="Close ($)"),

```

```

209         color=alt.Color("variable:N", title="Series"),
210         tooltip=[
211             alt.Tooltip("Date:T", format="%Y-%m-%d"),
212             alt.Tooltip("variable:N", title="Series"),
213             alt.Tooltip("value:Q", title="Close", format="$.2f"),
214         ],
215     )
216     .properties(height=300)
217 )
218
219 st.altair_chart(test_chart, use_container_width=True)
220
221 # Multi-step forecast
222 with st.spinner("Forecasting..."):
223     full_scaled = scaler.transform(values)
224     window = full_scaled[-look_back:].reshape(1, look_back, 1) # -look_back: go back
30 or 60 days until the end. Reshape it into 3 dimen.
225     fut_scaled = []
226     for _ in range(int(horizon)): # Loop through each forecast horizon, example,
14days.
227         # 1. Initially, it will from last data go back 30 or 60 days.
228         # 3. Recursive function, the window will be predict again
229         nxt = model.predict(window)
230         fut_scaled.append(nxt.item())
231         # 2. window will remove the last element and add the next prediction in the
window, and do the predict again.
232         window = np.concatenate(
233             [window[:, 1:, :], # will remove the oldest value, 1: slice the 1st
element, from 1..end
234             nxt.reshape(1, 1, 1)], # reshape and append the new predicted value
235             axis=1
236         )
237
238     fut_scaled = np.array(fut_scaled).reshape(-1, 1) # Hx1 scaled predictions
239     fut = scaler.inverse_transform(fut_scaled).ravel() # back to dollars
240
241     last_date = series.index[-1] # last timestamp in your selected data
242     forecast_idx = pd.date_range(
243         last_date + pd.Timedelta(days=1), # start the day *after* last_date
244         periods=int(horizon), # how many future timestamps
245         freq="B" # 'Business day' frequency (Mon-Fri)
246     )
247     forecast_df = pd.DataFrame({"Forecast_$": fut}, index=forecast_idx)
248
249     st.markdown("### Forecast")
250     c1, c2 = st.columns([2, 1])
251     with c1:
252         recent = series.iloc[-max(100, look_back + 1):].copy()
253         recent["Type"] = "Recent Close"
254         fc = forecast_df.copy()
255         fc["Type"] = "Forecast"
256         plot_df = pd.concat([
257             recent.rename(columns={"Close": "Close_$"}),
258             fc.rename(columns={"Forecast_$": "Close_$"})

```

```

259         ])
260
261     chart = (
262         alt.Chart(
263             plot_df.reset_index(names="Date")
264         )
265         .mark_line(point=True)
266         .encode(
267             x=alt.X("Date:T", axis=alt.Axis(format="%Y-%m-%d", title="Date")),
268             y=alt.Y("Close_$:Q", title="Close ($)"),
269             color=alt.Color("Type:N"),
270             tooltip=[
271                 alt.Tooltip("Date:T", format="%Y-%m-%d"),
272                 alt.Tooltip("Close_$:Q", title="Close", format="$.2f"), # <- use
Close_$
273                 alt.Tooltip("Type:N", title="Series"),
274             ],
275         )
276         .properties(height=300)
277     )
278
279     st.altair_chart(chart, use_container_width=True)
280
281     with c2:
282         st.dataframe(forecast_df.round(2))
283         csv = forecast_df.reset_index().rename(columns={"index":
"Date"}).to_csv(index=False).encode("utf-8")
284         st.download_button("Download forecast CSV", csv, file_name=f"
{ticker}_forecast.csv", mime="text/csv")
285
286         st.info(f"Last known date in range: {last_date.date()} – Forecast starts on
{(last_date + pd.Timedelta(days=1)).date()} for {int(horizon)} day(s).")
287
288     else:
289         st.info("Adjust parameters on the left, then click **Train & Forecast**.")

```